DHAR

Volume I — Vision & Specification

SPEC0

“One language, from source text to silicon, owned end to end.”

Document	Spec0 — Overview & Vision
Version	1.2
Date	September 7, 2026
Status	Ratified
Part of	The Canon of Dhar, Volume I
Author & Project Lead	Abdullahi Baba Garba

Philosophy & Naming

Dhar comes from the pre-Islamic Arabic *dahr*: time in its most unbounded sense — a long age, eternity, fate. Classical Arabic lexicons and the Quranic corpus confirm this as the word's core meaning. In pre-Islamic Arabian poetry, dahr was not merely a measurement of duration but an autonomous force — the thing that outlasts kings, wealth, and strength alike, answerable to nothing else. The concept was significant enough that Islamic tradition addresses it directly: a hadith recorded in both Sahih al-Bukhari (4826) and Sahih Muslim (2246) has Allah say, “the son of Adam curses time [dahr], though I am time” — read by classical commentators, al-Nawawi among them, as a correction of the pre-Islamic habit of blaming misfortune on time itself, not as a claim that time is divine. Dahr remains, in every source, time — total, autonomous, and enduring past any single moment or ruler.

That is the inheritance this project draws on, stated plainly rather than overclaimed: independence, robustness, and endurance are not dahr's dictionary definition, they are this Canon's chosen extension of it — the same way most technical names build meaning onto a word rather than translating it literally. A systems language and toolchain meant to outlast the hardware it targets and the individual engineers who build it aspires to the same quality dahr was always used to describe: something whole, self-standing, and built to last. Dahr's other defining quality — encompassing every fortune and misfortune, the entire span of a life rather than one register of it — is where this Canon locates Dhar's multi-paradigm ambition (§2.2 below): a language equally at home in a kernel driver and a web backend, in strict static typing and dynamic scripting, claims a similar totality.

Kun, the name given to Dhar's package manager (Stage-delta, Volume II), comes from the Arabic *kun*, “be” — the imperative at the center of the Quranic phrase *kun fa-yakun*, “be, and it is”, appearing eight times in the Quran (2:117; 3:47, 59; 6:73; 16:40; 19:35; 36:82; 40:68) to describe creation by divine command. A package manager's core promise has the same shape, at a vastly smaller scale: declare what should exist in a project, and it comes into being — resolved, fetched, and ready — without being assembled by hand. The metaphor is offered in that spirit, not as a claim of equivalence with what the phrase means in its own context.

1. Executive Summary

Dhar is a from-scratch systems programming language and compiler toolchain, designed around a single, deliberately ambitious target: performance at least at C's level, Java-class safety, Python-class readability *and* Python-class dynamism where it's wanted, genuine multi-paradigm support, and full independence from any external compiler backend or runtime. No LLVM. No JVM. No borrowed interpreter. The compiler is bootstrapped once and then produces its own machine code through its own backend, permanently.

Dhar's motivation is Nigerian and African technological independence at the systems-software layer — the layer beneath applications, where compilers, linkers, and runtimes live, and where almost no African-built alternative currently exists. This document is Spec0: the front door to the Canon. It states the problem, the technical thesis, the architecture, the target use cases, a taste of the language, and how Dhar is meant to sit relative to the languages that already exist. The full grammar and ratified keyword semantics live in Spec1. The stage-by-stage engineering plan lives in Volume II.

2. The Problem

2.1 Technological Dependence

Most of the software infrastructure the world runs on — compilers, operating-system kernels, virtual machines, cryptographic libraries, the toolchains that turn source code into running systems — was designed, and continues to be maintained, outside Africa. That is not a criticism of any of it: C, Rust, Java, Python, and the compilers behind them are mature tools built by large communities over decades. But it means that when African institutions build the software critical infrastructure runs on — banking cores, telecom switches, government platforms, defence systems — they build it on foundations they did not design, cannot unilaterally audit at the lowest level, and do not set the roadmap for.

Dhar starts from a simple observation: Nigeria, and Africa more broadly, does not currently have a home-grown, from-scratch systems programming language and toolchain — something built at the same layer as C, Rust, or Go, by people who can audit, extend, and maintain every layer of it locally. Closing that gap is not only about the language. It is about growing a generation of compiler engineers who understand their tools from the lexer up, building software that can be exported rather than only imported, and giving the institutions that most need to trust their infrastructure a toolchain whose internals are not a black box.

2.2 The Missing “Holy Grail”

The target combination this Canon builds toward is informally called the “Holy Grail” of programming languages: not a guarantee, but a target specific enough to force real trade-off decisions. No mainstream language currently holds all of the following at once:

- **C-class performance** — near-hardware, at least matching C, not merely approaching it: no GC pauses, no interpreter overhead, memory laid out and controlled at compile time.
- **Java-class robustness** — strict typing, with safety mechanisms such as automatic bounds-checking making silent memory corruption structurally difficult, not merely discouraged.
- **Python-class simplicity and dynamism** — code that reads close to plain language, with type inference doing the work a beginner would otherwise do by hand, and a dynamic-typing mode (Spec1 §4.4) available for exploratory, script-like code, the same way Python code can be written without ever declaring a type.
- **Genuine multi-paradigm support** — procedural, functional, object-oriented, logical, and symbolic styles combined intelligently in the same language (Spec1 §10), not bolted on.
- **Full independence** — no LLVM, no JVM, no external interpreter runtime, because the compiler is bootstrapped and produces its own machine code through its own backend.

Every existing language trades at least one of these away: C and C++ trade safety and simplicity for performance; Java trades independence and some performance for safety and portability; Python trades performance and low-level control for simplicity; Rust gets close on performance and safety but keeps a steep learning curve and an LLVM dependency; Zig and Odin get close on independence-in-spirit but still lean on LLVM for code generation on most targets today. Naming the full combination as the target is what lets every later decision in this Canon be checked against it.

> *Linux did not win because Unix was rare. It won because it was a complete, independently built, freely extensible foundation that the world could stand on without asking permission. Dhar's*

ambition at the language-and-toolchain layer is the same kind of foundation, not a niche experiment.

3. The Dhar Thesis

Five linked design commitments answer “how.” Each is expanded into full engineering detail in its corresponding Volume II stage document.

1. The parser is not written in assembly — the backend is.

Hand-writing a full lexer, parser, and type checker directly in assembly does not scale; parsing complexity would bottleneck progress before the language is usable. Assembly is reserved for what it's good at: a library of code-generation templates — pre-written structures with the hardcoded values stripped out — that a higher-level compiler fills in and emits.

2. A universal virtual CPU is defined before any real one is targeted.

The same way Python and Java compile to an intermediate bytecode rather than directly to hardware, Dhar defines its own universal instruction layout, not tied to any physical architecture. The compiler only ever targets this universal layout; a small backend module per architecture (x86-64, ARM64, RISC-V, WASI, Windows) turns it into real machine instructions. Portability is structural: a new architecture means writing one backend module, never touching the parser or type system.

3. Manual memory management is traded for compiler-inserted discipline.

Python's simplicity comes from not managing pointers by hand; C's speed comes from having no background garbage collector. Dhar keeps the manual-management-free surface syntax while the compiler automatically inserts allocation and deallocation logic during compilation, following a compile-time ownership model (Rust's approach) combined with defer/arena allocation (Zig's approach) where appropriate.

4. Safety is baked into the templates, not into a garbage collector.

Every array access is generated with an automatic bounds check, so an invalid index produces a safe, diagnosable crash instead of silent corruption — with no separate runtime safety layer required.

5. The language is bootstrapped through a disposable bridge, then proven by rebuilding itself.

A minimal bridge compiler, written in a language that already works, compiles just enough of Dhar to build a real compiler written entirely in Dhar. That compiler then compiles itself; if the resulting binary is byte-for-byte identical to the one that produced it (ignoring timestamps), Dhar has proven it no longer needs the bridge, which is retired to a historical artifact rather than kept as a permanent dependency.

4. Architecture Overview

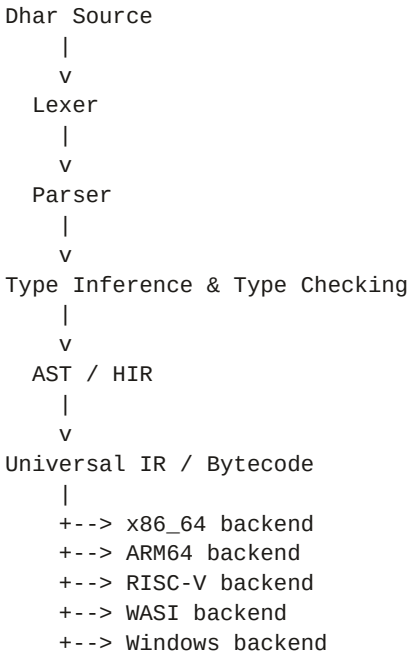

The load-bearing rule: nothing above the Universal IR line is allowed to know which hardware it is ultimately targeting. The lexer, parser, type checker, and AST construction only ever need to understand Dhar's own syntax and semantics. All architecture-specific knowledge is pushed into backend modules that consume the same IR and each know how to lower it to one target. Early code generation is specified to lean on the template-emitter approach from the Thesis; Stage-gamma (Optimizing Dhar) is where a real intermediate representation with optimization passes takes over.

5. Use Cases & Target Domains

- **Backend & server software** — C-class performance without a GC pause.
- **Web** — compiled backends, with WASI as a first-class backend target.
- **Mobile** — via the ARM64 backend.
- **Cybersecurity & security tooling** — memory-safety guarantees plus low-level access (`peek`/`poke`, `raw[N]`) in one language.
- **Data & AI workloads** — performance-critical numeric code without leaving a readable syntax.
- **Blockchain** — deterministic execution and fine-grained resource control.
- **General systems programming** — OS components, drivers, and other code currently written in C or Rust.
- **Hardware & embedded programming** — direct memory operations and fixed-size buffers as language primitives.
- **Defence & critical national infrastructure** — a toolchain whose entire pipeline, from lexer to machine code, can be independently audited end to end, with no externally controlled dependency in the critical path. Specified

in full in Tempering Dhar (Volume III).

6. Language Preview

Dhar defaults to Python-like concision, with explicit typing available for systems-level control — both forms are first-class and coexist in the same program.

```
# concise, inferred -- unqualified bindings are mutable by default
name = "Dhar"
count = 10
count = count + 1        # fine: count was never locked

# explicit, for systems-level code or to opt into immutability
lock pi: f64 = 3.14159
flux total: i32 = 0
address: raw[4096]

# collections, inferred (Spec1 S11)
scores = [98, 87, 92]              # a list
lock bonus = (n: i32) -> n + 5     # a closure -- functions are first-class
each s in scores:
    gen(bonus(s))

# structs and methods -- brace-delimited, unlike the rest of the language
mold Point {
    flux x: i32
    flux y: i32
}
forge Point {
    task init(a: i32, b: i32):
        x = a
        y = b
}
task dist(p: view Point): i32     # view = read-only reference
    give p.x + p.y

# modules, entry point, and error handling
pull math
expose dist
task core():                      # core() is the standard entry point
    flux s: state = may_fail()
    trap s
    enforce s == 0
```

This is a taste, not the specification. Every keyword shown above — the full 46-keyword set (including generics, sum types, interfaces, and async), the type system, the memory model, and collections — are ratified in Spec1.

7. Positioning: How Dhar Compares

Language	Performance	Memory Safety	Toolchain Independence	Multi-Paradigm Depth
C	Near-hardware	Manual, unchecked	Depends on GCC/Clang	Procedural only
C++	Near-hardware	Manual + optional RAII	Depends on GCC/Clang/MSVC	Procedural + OOP + generics
Rust	Near-C	Compile-time ownership	Compiler self-hosted; codegen depends on LLVM	Mostly procedural/OOP
Java	JIT, GC pauses	Managed (GC)	Depends on the JVM	OOP-centric
Python	Interpreted	Managed (GC)	Depends on the CPython runtime	Dynamic, not symbolic-native
Go	Compiled, GC pauses	Managed (GC)	Self-hosted, own backend	Procedural/OOP-lite
Zig	Near-C	Manual + compiler checks	Frontend independent; codegen still leans on LLVM	Procedural
Dhar	Near-hardware — at least C-class	Ownership/arena + template-level bounds checking	Fully self-hosted, custom IR, no LLVM/JVM	Procedural, functional, OOP, logical, symbolic — plus optional dynamic typing

Dhar's row states this Canon's design commitments, deliberately placed in the same performance tier as C rather than merely "close to" it. Stage-gamma (Optimizing Dhar) carries the performance commitment; Spec1 carries the safety and paradigm commitments; Stage-alpha and Stage-beta carry the independence commitment.

8. The Five Stages

Stage	Name	Mandate
Stage-alpha	Forge of Dhar	Bootstrap compiler proving the pipeline end-to-end for a minimal language subset
Stage-beta	Honing Dhar	A complete compiler written entirely in Dhar, capable of compiling itself
Stage-gamma	Optimizing Dhar	SSA, register allocation, jump tables, vectorization
Stage-delta	Gearing Dhar	Standard library and tooling: build, test, format, packages, LSP
Stage-epsilon	Deploying Dhar	Cross-compilation, debug info, fuzzing, reproducible builds, SBOM, signing

Stage-gamma must not begin until Stage-beta is genuinely self-hosting and the syntax ratified in Spec1 is stable. Full mandates, acceptance criteria, and required tests for each stage are specified in their own Volume II documents.

9. Reading the Rest of the Canon

- Writing Dhar code or building the compiler → Spec1, then the relevant Volume II stage document.
- Evaluating Dhar for adoption, funding, or partnership → Casting Dhar and Fueling Dhar.

- Assessing Dhar for security, critical-infrastructure, or defence use → Tempering Dhar, alongside Spec1.
- Joining as a contributor → The Guild of Dhar.
- Tracking overall progress → The Pact of Developing Dhar.

The full index lives in The Canon of Dhar — Master Index. If Linux is the proof that an independently built, freely extensible foundation can outgrow the proprietary systems around it, this Canon is Dhar's plan for the same outcome one layer up the stack, at the language itself.

Spec0 — Overview & Vision · v1.2 · September 7, 2026 · Ratified. Revision log: v1.0 — rewritten as a pure forward-looking specification. v1.1 — fixed mold/forge to brace-delimited syntax for consistency with Spec1; unqualified bindings now documented as mutable (flux) by default; performance positioning restated as near-hardware/C-class rather than merely near-C; keyword count updated to 38 following the addition of wild (Spec1 §4.4); added the Linux analogy. v1.2 — added §0, Philosophy & Naming, grounding Dhar and Kun in researched etymology; keyword count updated to 46 following Spec1 v3.0.